

Lab 3

Matt Asaro
Lawrence Lee
Charles McKay

Preliminaries

1. *Perform the following tasks for your non-pipelined processor and include them in your written report:*
 - a. *List the reported Fmax of your design before pipelining*
52.11 MHz
 - b. *List the cycle time of your design before pipelining. Show your work*
 $1/52.11 \text{ MHz} = 1/5211000 \text{ Hz} = \mathbf{19.2 \text{ nanoseconds}}$
 - c. *List the instruction count for miner_tb from the ModelSim simulation before pipelining.*
120561 instructions
 - d. *List the cycle count for miner_tb from the ModelSim simulation before pipelining.*
131490 cycles
 - e. *Compute the CPI and execution time (# of cycles x cycle time) before pipelining.*
ET = 2.52 ms
CPI = 1.09
 - f. *List the number of registers used, the number of combinational functions used, and the number of memory bits used before pipelining.*
2069 registers
4419 combinational functions
16384 memory bits
 - g. *Include a datapath diagram of your core. We strongly encourage you to draw this by hand. It will become a very useful tool in debugging your core.*
2. *Include answers to the following questions in your written report:*
 - a. *On your datapath diagram, define the boundaries of stages in your pipeline by drawing in the required registers. Spend some time doing this*

The diagram illustrates a MIPS-like processor architecture with the following components and connections:

- PC (Program Counter):** Outputs an address to **I.mem**.
- I.mem (Instruction Memory):** Outputs an instruction to a multiplexer (labeled 1).
- C1-decode:** Receives the instruction and outputs control signals: *is_load*, *write_reg*, *is_store*, *is_mem*, and *is_byte*.
- Reg File (Register File):** Receives the instruction and outputs register addresses (*rs_addr*, *rd_addr*, *wr_addr*, *wr_data*) and register values (*rs*, *rd*). It also receives control signals *is_load* and *write_reg*.
- ALU (Arithmetic Logic Unit):** Takes register values and performs operations based on control signals. It outputs a result to the **Data Mem**.
- Data Mem (Data Memory):** Receives data from the ALU and outputs data to a multiplexer (labeled 3) and back to the Reg File. It also receives control signals *is_store*, *is_byte*, and *is_mem*.
- Control Logic:** Includes several multiplexers and logic blocks:
 - Multiplexer 1: Selects the instruction from I.mem.
 - Multiplexer 2: Selects the next PC value from the ALU result or other sources.
 - Multiplexer 3: Selects the data to be written back to the Reg File.
 - 5-bit equality and 6-bit equality blocks: Compare register values for branch instructions.

- IF
No control signals

EX/Mem
is_store_op
is_byte_op
is_mem_op
is_load_op (for stage 4)
op writes rf (for stage 4)

```
is_load_op
op_writes_rf
```

- [illegible]

- a. List all data hazards in your pipeline, and how they should be resolved (which stage data should be forwarded from, etc).

We will use pipe cuts to forward all data values through pipe cuts until they are no longer needed. Specifically, we will compare the `w_addr` in stage 4 to stage 3 `rd_addr` and `rs_addr`. If either values match, then we will forward and select the `rf_wd` value instead of the respective stage 3 values.

- ```
is_store_op (stage 3), is_byte_op (stage 3), is_mem_op (stage 3),
is_load_op stage 4), op_writes_rf (stage 4)
```

**Control hazards are resolved by forwarding control signal values through pipe cuts until they are no longer needed.**

- c. List all structural hazards in your pipeline, and how they should be resolved.

d. JALR, BEQZ, BNEQZ, BGTZ, BLTZ

**We will stall new instructions from writing to pipe cuts during the execution of these instructions.**

5. *Why would you want to have an independent write address port on the register file for a pipelined processor (i.e. why would you need to write to an address other than the rd address of the instruction in the decode stage)? Hint: think about what instructions may be in the write back stage (the stage which writes to the register file) and the the decode stage (the stage that reads from the register file).*

**Because in a pipelined processor the write data is computed in a different cycle than the read. When the write data is ready to be written to the register file the processor is already decoding another instruction which can have different read and write addresses.**

Tasks completed:

- The first thing we noticed was that none of our instructions for accessing data memory had any offset. It appeared that the ALU was acting as a mux for the data memory input values. As a result, the ALU and Data Memory will happen in the same stage, in parallel, since the Data Memory is independent of the ALU. This means that we can max out at 3 pipe cuts (4 stages) instead of what is recommended in the lab write up.
- We flush out all pipecuts if we get a reset signal.
- The first pipecut collects the instruction fetched and the PC count of that instruction. It needs the PC for jump/branch calculation. In this first stage, we insert nops if we encounter a jump, branch, or wait instruction. We got this pipecut working 100%.
- We put in the second pipe cut forwarding the instruction, pc, rs and rd address and values. If we calculate a value for a register to be used in the next instruction we write that value into the second pipecut rather than the fetched value of the register. That value is stale. **This is our first use of data forwarding.** However, we currently don't pass the miner test as it finds it on cycle 22.

Tasks Left:

- Finish debugging of Cut 2
- Add Cut 3
- Jump buffer register
- branch prediction
- miner.asm optimizations


## Optimization Results

1. Include answers to the following questions in your written report:

a. How many stages does your pipelined processor have?

**4**

b. Which pipeline cuts did you successfully implement? Include a diagram.



2. Include answers to the following questions in your written report:

a. List the reported Fmax of your design after pipelining. **126.55 MHz**

b. List the cycle time of your design after pipelining. Show your work.

$$1/(126.55 \text{ MHz}) = 7.9 \text{ ns}$$

c. List the instruction count for miner\_tb from the ModelSim simulation after pipelining.

**98279**

d. List the cycle count for miner\_tb from the ModelSim simulation after pipelining.

**109198**

e. Compute the CPI and execution time (# of cycles x cycle time) after pipelining.

$$ET = 109198 * 7.9 \text{ ns} = 826.7 \text{ } \mu\text{s}$$

$$CPI = CC / IC = 109198 / 98279 = 1.1111$$

f. List the number of registers used, the number of combinational functions used, and the number of memory bits used before pipelining.

**Registers: 296**

**Combinational Functions: 1,481**

**Memory Bits: 20,480**

g. What is your speedup (baseline execution time/optimized execution time)?

$$\text{Speedup} = 2.52 \text{ ms} / 826.7 \text{ } \mu\text{s} = 2.92$$

- h. How does pipelining differ from the optimization method used in Lab 2 (adding instructions to the ISA)? Specifically which of the two methods targets cycle time, instruction count, CPI, etc. Hint: You may want to compare your results from Lab 2 with your results from Lab 3.
- Pipelining increases throughput by decreasing cycle time, but also increases CPI, whereas adding instructions decreases the instruction count and increases CT (at least the way we did it in Lab 2).**

## Additional Notes/Explanations

### ALU/Data Mem

The pipeline cuts we implemented were IF/ID, ID/(EX, MEM), and (EX, MEM)/WB. The first task we accomplished was putting the ALU and the data memory in parallel. Originally, the ALU was the third stage (EX), and the data memory was in the fourth (MEM). We saw that there are no memory instructions that depend on the ALU, as offsets do not need calculation. In fact, the ALU was simply acting as a mux for various inputs into the data memory.

So we moved the data memory to the EX stage, making it the (EX, MEM) stage, and we also added a mux that performed the same “mux” functions the ALU did prior. This ended up decreasing clock time and the total number of clock cycles, as the instruction latency became 4 instead of 5.

As such, our CPU has 3 pipecuts because there only are 4 stages.

### First Pipecut: IF/ID

The first pipecut consists of two registers: one for the instruction fetched and one for the PC. The main difficulty with the first pipecut was figuring out that the PC of the fetched instruction needs to be saved and forwarded, as branch and jump address calculations depend on the PC that fetched the branch or jump instruction.

In addition, branches and jumps resolve in EX by default, so we put in checks that flush/stall when necessary.

### **Second Pipecut: ID/(EX, MEM)**

The second pipecut has the most saved registers. In addition to the instruction and the PC, it contains all five control signals and the values from the register file. We had to continue various checks for branch resolution, add register file forwarding, and add data forwarding from the ALU and data memory back to this pipecut.

Our issue with implementing this pipecut for the milestone was that we did not properly have data forwarding.

### **Third Pipecut: (EX, MEM)/WB**

Having resolved all branches and jumps, we did not have to do nearly as much in this pipecut as we did in the previous one. Only the instruction, the address/value for writing back into the register file, and various control signals needed to be put into this pipecut.

### **Data Forwarding Details**

We implemented data forwarding at two points in the CPU:

1. Between the write address in (EX, MEM) and rd/rs addresses in ID
2. Between the write address in WB and rd/rs addresses in ID (register file forwarding)

We checked for equality between the addresses at those points and, if necessary, wrote the correct values from either (EX, MEM) or WB into the second pipecut (ID/(EX, MEM)) instead of the values read from the register file. This data forwarding eliminates any stalls that we would have had to insert because of data dependencies.

### **Jump Optimizations**

We moved jump address calculation from EX to ID in order to stall for fewer cycles on a jump. We did this by simply changing the calculations to use the data from the instruction in the first pipecut (IF/ID) instead of the instruction in the second pipecut (ID/(EX, MEM)). This doubled the CPI of the JALR instruction.

We also optimized miner.asm by removing unnecessary jumps, which reduced stalls and brought our overall CPI down. **Note that without the miner.asm optimizations, we still achieved a speedup of over 2.**

### **ALU Optimizations**

We removed processing for instructions that are unused in miner.asm, which significantly increased Fmax. For example, kOR simply returns 0, as OR is never used in miner.asm.